

# Keith R. Bennett's Technical Blog

About

## Design by Contract, Ruby Style

By [keithrbennett](#) on June 15th, 2011

I recently encountered a situation in which I was writing data to a key/value data store. There was a code layer that insulated the business logic layer from the data store internals. I wanted to be able to unit test this business logic without needing access to the data store.

I could mock the data access layer, but I wanted something more functional — something that would *behave* like the data store layer, and possibly even be used to test it. I decided to write something mimicking the production data store layer that used a Ruby Hash for storage. How could I use the language to help me know that I had faithfully reproduced all the functions in the original object?

## Dependency Inversion

Dependency inversion is defined in [Wikipedia](#) as:

- A. High-level modules should not depend on low-level modules. Both should depend on abstractions.
- B. Abstractions should not depend upon details. Details should depend upon abstractions.

Using this principle, I rearranged the business layer so all it needed to know about its data access object was the method signatures and nothing else, not even its type. The data access object was passed into the business object's constructor rather than the business object instantiating it directly (and thereby needing to know its type). That opened the door to alternate data access strategies, including a trivial implementation for unit testing.

## Java Interfaces

If you've coded in Java, you'll recognize that this principle is what Java interfaces are all about. Using Java interfaces involves the following:

- 1) defining the interface, that is, the signatures of the methods that need to be defined by implementations of that interface.
- 2) defining object references in the code that refer to the implementations as the interface type rather than the implementation type.
- 3) creating implementations of that interface.

## Ruby's Duck Typing

Ruby uses duck typing, which eliminates the need for #1. Unlike Java with its compile time checking, Ruby doesn't verify the existence of the method before it's actually called — it trusts the method to be there, and calls it, raising an error if it doesn't exist, and is not handled by a `method_missing` function. In addition, in Ruby, references are not typed, eliminating the need for #2.

In this way, Ruby makes dependency inversion effortless and automatic.

Many Ruby developers believe the Java interface approach to be obtrusive and verbose. However, there can be issues with Ruby's permissive approach as well. For example, while relying on duck typing is convenient, we want noncompliant implementations to fail quickly and gracefully. Using duck typing alone, the absence of a method will be detected only when it is called — and in the absence of adequate unit testing (which could be beyond one's control), this could result in premature termination of a long job, or the detection of the error months after its introduction.

In addition, it would be helpful to show the user *all* the methods that are missing, not just the one that was called. Duck typing will not do that; it will throw a `NoMethodError` only for the method called.

There are other problems with relying exclusively on duck typing. Where can the author of a new implementation of the contract find all the methods that need to be implemented? Maybe one of the implementations? Which one? And can we trust that the author of that implementation made private all the methods that are not needed in the public interface?

## Making the Promise in Executable Code

For the same reason that self documenting code is better than comments, these required methods should be expressed as executable code rather

than included in some documentation file. Further, this promise is such an important part of the class' existence, that it makes sense to include it somewhere inside the class definition. (The promise could live here even if we only validate the class' adherence to the contract during unit testing.)

Sometimes this will be overkill — if we have complete control over our source code, then we can unit test to our heart's delight, and guarantee that only thoroughly tested code makes it to production. However, there are some situations in which more verification would be helpful.

## An Illustrative Example

Let's say we're writing a flexible data processing framework that allows users to plug in their own implementations of actions. As a trivial example, let's say that one of the services of this framework is encrypting strings, and we want the users to be able to provide their own encryption strategies.

Let's assume there are lots of these plugins used by the system, some provided by the software authors and others provided by customers and third parties. Let's also assume that we want to verify that these plugins implement the required methods before starting the job. (The job may take a long time, use expensive resources, be especially critical, etc.)

What is the appropriate Ruby way to provide some kind of confidence in these plugins? Certainly, we should be providing implementation authors with comprehensive test suites that test much more than the presence of required methods. However, let's go further than that. Let's provide a mechanism to enable this verification at any time, not just in unit tests.

Here is a class that can be used to verify the presence of instance methods (also at <https://gist.github.com/1034775>):

```
1  # Shows a crude way to enforce that a class implements
2  # required methods. Use it like this:
3  #
4  # class OkClass
5  #   def foo; end
6  #   def bar; end
7  #   def baz; end
8  # end
9  #
10 #
11 # class NotOkClass
12 #   def foo; end
13 # end
14 #
15 #
16 # verifier = InstanceMethodVerifier.new([:foo, :bar, :baz])
17 # verifier.verify('OkClass')
18 # verifier.verify('NotOkClass')
19 #
20 # $ ruby instance_method_verifier.rb
```

```

21 # instance_method_verifier.rb:22:in `verify': Cl
22 #   from instance_method_verifier.rb:42
23
24 class InstanceMethodValidator
25
26   attr_accessor :required_function_names
27
28
29   def initialize(required_function_names)
30     @required_function_names = required_function_
31   end
32
33
34   # klass can be either a class name or a class o
35   def validate(klass)
36
37     if klass.is_a? String
38       klass = Kernel.const_get(klass)
39     end
40
41     missing_function_names = required_function_nc
42
43     unless missing_function_names.empty?
44       raise Error.new("Class #{klass} is missing
45     end
46   end
47
48
49   class Error < RuntimeError
50   end
51 end

```

The contract below (also at <https://gist.github.com/1034784>) lists the required methods. It's also a good place to provide information about these methods.

```

1  require 'instance_method_validator'
2
3  module EncryptionContract
4
5    def required_function_names
6      [
7        # The purpose of this function is self-evident.
8        # cases the methods would benefit from some
9        # that would be helpful to authors of new i
10       # such as defining correct behavior, format
11       # of arguments and return values.
12       #
13       # encrypt a string
14       #
15       # arg: the string to encrypt
16       # returns: the encrypted string
17       'encrypt',
18
19       # decrypts a string
20       #
21       # arg: the string to decrypt
22       # returns: the decrypted string
23       'decrypt', # (encrpyted string)
24     ]
25   end
26
27
28   def validate_contract(klass = self.class)
29     InstanceMethodValidator.new(required_function_
30   end
31 end

```

Now let's write a naive example implementation. Having the above code makes it trivial for the implementer to test that all required methods are implemented. Here is the rspec (also at <https://gist.github.com/1034785>):

```

1 | require 'rspec'
2 | require 'reverse_encrypter'
3 |
4 | describe ReverseEncrypter do
5 |
6 |   it "should implement required methods" do
7 |     lambda { subject.validate_contract }.should_not
8 |     end
9 |   end

```

Here is the implementation itself (see also

<https://gist.github.com/1034788>):

```

1 | require 'encryption_contract'
2 |
3 | # Naive and totally lame encrypter that just reverses
4 | class ReverseEncrypter
5 |
6 |   include EncryptionContract
7 |
8 |   def encrypt(s)
9 |     s.reverse
10 |   end
11 |
12 |   def decrypt(s)
13 |     s.reverse
14 |   end
15 | end

```

In a pluggable framework as described above, it can be helpful to validate all “foreign” (that is, not provided by the framework authors) components as thoroughly as possible. The approach above enables the validation to be done at will, either in unit tests or at runtime. It's a trivial function call and takes virtually no time to execute.

Another benefit of using this approach is that it naturally results in more thought given to the public interface. Do we really need all those methods? Do I need to change any method names so they don't hint at implementation? Do I need to use a higher level object in my calling code to better insulate myself from the implementation?

## Conclusion

This approach is no substitute for thorough unit testing. On the contrary, it is an *aid* to unit testing, and more.

Omitting a function is just one of an infinite number of ways to screw up an implementing class. This approach doesn't verify that the arguments to the function are consistent across implementations, and that is very important too. Nor does it help test behavior. Nevertheless, given that limitation, it has the following benefits. It:

- formalizes the contract in executable code, where it is far more likely to be current
- provides an authoritative source of that contract
- promotes thinking more about the public interface
- aids in writing a new implementation class
- makes testing implementations DRYer; required method names are in one place rather than multiple unit tests, one for each implementation, so tests are more likely to be present, and more likely to be correct
- can be validated outside of unit tests

I think that's pretty worthwhile. That said, I'd be interested in better ways. Feel free to comment on this article.

– Keith

Categorized under: [Uncategorized](#).

Tagged with: no tags.

## 2 Responses to “Design by Contract, Ruby Style”



**Russ Olsen** says:

June 16, 2011 at 8:16 am

Keith,

First let me say that I applaud your efforts – it's exactly this sort of 'what about this...' thinking that helps everyone get better. I do think that your ideas show how hard it is to shift modes of thinking,

If I understand your code correctly, you are building a class that verifies that some other class implements a given set of instance methods in a effort to avoid typing problems at runtime. Here are some things that I think you should consider:

\* A Ruby class does not completely describe the behavior of it's instances. Instances are free to define singleton methods of their own, pull in modules etc. This is not terribly common, but it is perfectly acceptable Ruby.

\* The methods returned via `instance_methods` do not necessarily describe all of the messages that instances of the class can respond to. Using `method_missing`, a class can build all sorts of behavior that `instance_methods` will miss.

\*This is a nit, but I think you should be using `public_instance_methods` rather than `instance_methods`, so that you don't accidentally see private methods which would be un-callable.

\* Also I think you could implement your 'is the method there?' code a bit cleaner with array arithmetic – it turns out that `(array1 - array2)` will give you an array with all of the elements in the first array that are not in the 2nd.

These are all details, however. The real issue is that years of working with static typing has left us all with the focus what methods a class supports and perhaps the types of their arguments and that sort of thing. Certainly in a language like Ruby you can have problems with methods not being there. But the 'this object does not support the right methods' bug is only one of an infinite number of mistakes that you can make. The crazy thing is that the type related issues that statically typed languages devote so much time and energy to are simply not that common compared to some of the other bone headed things we do everyday. As others have pointed out on the novarug mailing list, the only way to hunt down those bugs is to write unit tests.

Think about the kind of test support that rails give you – I can't think of any magic rails thing that checks to see if this class has the right methods. But what you do have is lots and lots of tools that help you test your objects in situ.

Anyway, those are my \$0.02. Thanks for putting this out there – no matter if I agree with your approach or not, this kind of carefully written article is always helpful.

Russ

[Log in to Reply](#)



*keithrbennett* says:

June 19, 2011 at 4:41 pm

Russ –

Thanks for visiting my blog. You make some good suggestions that I will incorporate into the article and my code (e.g. array subtraction and use of the `public_instance_methods` method).

Among the languages I've learned over the years, Java is the most formal and Ruby the most informal. From where I stand, my strategy has nothing to do with an incomplete grasp of the Ruby way; rather, in my opinion, my use case requires something a little

more formal than Ruby duck typing alone. We are (hopefully) software developers and problem solvers first, and Ruby developers second — the operative questions should be “does this make sense?”, “do the benefits exceed the costs?”, etc., more than “do other people do it like this in Ruby?”. While I’m not saying that the latter question is devoid of merit, its weight should be tempered by the nature of the problem at hand. I realize that your reasoning goes well beyond adherence to the Ruby way, but it sounded to me like that was part of it.

What you say about `method_missing`, singleton methods, modules, etc., is true. However, again, I am making an architectural decision to formalize the contract — in that context, I think it’s reasonable to require that a function be explicitly defined. In my case, these classes are written specifically for this framework, so it’s not likely that people have classes lying around that adhere to the contract but use the other mechanisms you mentioned; and even if they did, a delegating method would be trivial to add.

As for the tools provided by Rails, I’m using Ruby but not Rails. Nor am I very familiar with Rails, so if it contains tools that would be helpful here, I would appreciate any pointers.

You’re not alone in your critique of my idea. Many others have said that unit tests are the *\*only\** place for this verification. However, in my opinion, there are some pretty compelling benefits I mentioned in the article which have not been addressed (see the bullet points at the conclusion for a summary).

By the way, I’ve renamed the article from “Dependency Inversion, Ruby Style” to “Design by Contract, Ruby Style”. Dependency inversion, Ruby style is provided automatically due to duck typing, so it wasn’t an appropriate title.

– Keith

[Log in to Reply](#)

## Leave a Response

You must be [logged in](#) to post a comment.



## ← Android GUI's — The Case Against GUI Builders and Data Driven GUI Configuration

class\_eval, instance\_eval, eval →

### META

- Register
- Log in
- Entries RSS
- Comments RSS
- WordPress.org


### ARCHIVES

- November 2015
- November 2013
- August 2013
- January 2013
- December 2012
- November 2012
- September 2012
- July 2012
- January 2012
- June 2011
- May 2011
- March 2010
- July 2009
- March 2009
- February 2009
- November 2008

### RECENT ACTIVITY

**Posts** **Comments**

- The Case for Nested Methods in Ruby
- Ruby's inject/reduce and

- each\_with\_object
-  in Your System Prompt
  - Using Oracle in JRuby with Rails and Sequel
  - Copying (RVM) Data Between Hosts Using ssh, scp, and netcat
  - Building A Great Ruby Development Environment and Desktop with Linux Mint 13 "Maya" Mate
  - Intro to Functional Programming in Ruby
  - WordPress Administration with Ruby
  - Stealth Conditionals in Ruby
  - Hello, Nailgun; Goodbye, JVM Startup Delays

ARCHIVES

Categories	Tags	Dates
Authors		
■ Uncategorized		